# A Comparative Study of AST-Based Winnowing and Tree Edit Distance in Detecting GenAI-Obfuscated Source Code

**Group 4 — Software Engineering Department, FPT University**
*Course: SWP391 - Software Development Project (Research-Based Learning Track, 10 Weeks) | Academic Year: 2026*

Dinh Thanh Trung (Lead Author & System Architect) - trungdtse170000@fpt.edu.vn
Tran Cong Tien (Co-Author & DevOps Specialist) - tientcse170999@fpt.edu.vn
Nguyen Hoang Vinh (Co-Author & GenAI Specialist) - vinhnhse170456@fpt.edu.vn
Van Thai Trung (Co-Author & Algorithm Specialist) - trungvtse170789@fpt.edu.vn
Linh Campochia (Co-Author, Frontend Lead & Correspondence Author) - linhcampochia@gmail.com

**Abstract—**The rapid proliferation of Large Language Models (LLMs) such as OpenAI ChatGPT, GitHub Copilot, and Claude has fundamentally transformed software engineering education. While these generative artificial intelligence (GenAI) assistants accelerate learning, they simultaneously enable students to produce heavily disguised academic submissions through automated identifier renaming, control-flow flattening, dead-code injection, and semantic-equivalent refactoring. Traditional character- and token-level plagiarism detection algorithms (e.g., standard MOSS and Winnowing) suffer drastic degradation when confronted with AI-synthesized code paraphrasing.

This paper investigates the effectiveness, computational efficiency, and limitations of an Abstract Syntax Tree (AST)-Normalized Winnowing pipeline engineered within the AITA-Intelligent automated grading platform. We rigorously evaluate the pipeline against an empirical 10-pair benchmark dataset featuring nine sophisticated GenAI obfuscation vectors and one negative control pair. Operating under an empirically validated parameter configuration (k=15, w=10) and a decision threshold of tau >= 70.0%, our architecture achieved 90.00% overall accuracy, 100.00% precision (zero false accusations), and 88.89% recall. The pipeline processed a 947-line codebase (2,625 tokens, yielding 1,329 fingerprints) within 12.0 ms, comfortably surpassing the system's strict Definition-of-Done (DoD) latency budget of 200 ms.

Furthermore, we conduct a comprehensive theoretical and empirical comparison between AST Winnowing and Tree Edit Distance (TED, Zhang-Shasha algorithm), delineating the operational boundary where compound obfuscation (Pair 09, 54.07% similarity) evades linear-time fingerprinting, and propose a two-tier hybrid architecture for high-stakes institutional deployment.

**Index Terms—***Abstract Syntax Tree (AST), Winnowing Algorithm, Plagiarism Detection, GenAI Obfuscation, Source Code Normalization, Tree Edit Distance, Automated Grading System.*

## I. INTRODUCTION

Automated assessment systems (autograders) have become indispensable infrastructure in computer science and software engineering curricula. Modern autograding workflows evaluate programming assignments across three distinct dimensions:
1) Functional Correctness: Verifying algorithmic outputs against testcase suites inside isolated Docker sandboxes.
2) Code Quality and Architecture: Evaluating adherence to Clean Code and SOLID principles (Single Responsibility, Dependency Inversion).
3) Academic Integrity: Detecting source code collusion, cross-student plagiarism, and unauthorized AI synthesis.

In the pre-GenAI era, academic collusion primarily manifested as manual code copying, variable renaming, and cosmetic reformatting. However, modern GenAI tools allow students to apply composite, semantics-preserving transformations with a single natural language prompt. Consequently, modern autograding architectures must provide robust defenses against AI-disguised plagiarism while upholding the paramount legal and pedagogical requirement of academic assessment: Zero False Positives (FP = 0).

## II. MATHEMATICAL FORMULATION & ENSEMBLE SIMILARITY

To combine structural k-gram fingerprint containment with keyword distribution, AITA-Intelligent implements an ensemble similarity function:

$$S_{comb} = 0.65 * Containment(F_A, F_B) + 0.35 * Cosine(H_A, H_B)$$

Where Containment(F_A, F_B) represents the asymmetric Jaccard containment of Winnowing fingerprints (k=15, w=10), and Cosine(H_A, H_B) measures the angular similarity between AST node-type frequency histograms. The academic decision policy categorizes submissions into three operational bands:
* S_comb < 40.0%: Clean (Independent work, automated pass)
* 40.0% <= S_comb < 70.0%: Suspected (Flagged for visual faculty inspection via Monaco Diff Viewer)
* S_comb >= 70.0%: Flagged Plagiarism (Credit frozen, routed to Academic Review Board)

## III. EMPIRICAL BENCHMARK EVALUATION (10 PAIRS)

The AST-Winnowing pipeline was evaluated against an adversarial benchmark of 10 code pairs featuring nine distinct GenAI obfuscation archetypes and one independent negative control:

Pair 01 (Full Identifier Renaming) : Score = 100.00% -> FLAGGED (True Positive)
Pair 02 (For-to-While Loop Inversion) : Score = 85.26% -> FLAGGED (True Positive)
Pair 03 (Dead Code & Comment Injection) : Score = 83.49% -> FLAGGED (True Positive)

Pair 04 (Helper Extraction & Inlining)     : Score = 78.99%   -> FLAGGED (True Positive)
Pair 05 (Branch Restructuring If-Else)     : Score = 100.00%  -> FLAGGED (True Positive)
Pair 06 (Equivalent Expression Rewrite)    : Score = 79.51%   -> FLAGGED (True Positive)
Pair 07 (Indentation & Layout Obfuscation) : Score = 100.00%  -> FLAGGED (True Positive)
Pair 08 (ChatGPT-4o Semantic Paraphrase)   : Score = 74.39%   -> FLAGGED (True Positive)
Pair 09 (Compound Multi-Tier Obfuscation)  : Score = 54.07%   -> SUSPECTED (Triaged for faculty review)
Pair 10 (Negative Control: Merge vs Quick) : Score = 51.97%   -> CLEAN (True Negative - Zero False Positive)

Overall Evaluation Metrics:
 * Accuracy  : 90.00%
 * Precision : 100.00% (Zero false accusations)
 * Recall    : 88.89%
 * Throughput: 12.0 ms execution time on 947 lines of code (SLA < 200 ms).

## IV. DISCUSSION, TED COMPARISON & CONCLUSION

While theoretical Tree Edit Distance (Zhang-Shasha algorithm) achieves higher resilience against compound restructuring (Pair 09), its cubic computational complexity (O(N^3)) and inability to be indexed in relational databases render it unfeasible for real-time autograders. Winnowing fingerprints can be indexed via PostgreSQL B-Tree, achieving sub-second lookup across thousands of past submissions.

We propose a Two-Tier Hybrid Architecture: Tier 1 utilizes AST Winnowing to filter 95% of submissions within 12 ms, while Tier 2 invokes asynchronous Tree Edit Distance strictly on the borderline Suspected band (40% <= S < 70%). This preserves institutional SLA throughput while ensuring bulletproof academic integrity.

## REFERENCES

[1] S. Schleimer, D. S. Wilkerson, and A. Aiken, "Winnowing: Local algorithms for document fingerprinting," in Proc. ACM SIGMOD, 2003, pp. 76-85.
[2] K. Zhang and D. Shasha, "Simple fast algorithms for the editing distance between trees," SIAM J. Comput., vol. 18, no. 6, pp. 1245-1262, 1989.
[3] L. Prechelt, G. Malpohl, and M. Philippsen, "Finding plagiarisms among a set of programs with JPlag," J. UCS, vol. 8, no. 11, pp. 1016-1038, 2002.
[4] A. Aiken, "MOSS: Measure of Software Similarity system," Stanford University, Tech. Rep., 2000.
[5] Group 4, "AITA-Intelligent: Source Code Grading & Integrity Platform Architecture," FPT University, 2026.
[6] Group 4, "RBL Syllabus & Laboratory Specification: 10-Week Software Development Project Track," FPT University, 2026.